

ELECTRICAL ENGINEERING
COMPUTER SYSTEMS ENGINEERING
ELE335 – COURSE PROJECT (TERM 2)



Project 1 & 2
Course: ELE335 (Artificial Intelligence)

S#	Student Name	Student ID
1	Mohamed Alaa	30
2	Abdelrahman Osama	13
3	Mohamed Hamdy	27
4	Abdullah Khaled	18
5	Mohamed Abdelmonem	29

Marks		
Free Topic Description		/40
Problem solving as a search problem	Problem formulation	/40
	Implementation/Evaluation /Sample output	/60
	Project interface & Poster Design	/25
	Report Style and Formatting	/15
	Presentation / response to questions	/20
	Creativity (Bonus)	/10
Expert System in Prolog	Implementation/Evaluation /Sample output	/60
	Project interface & Poster Design	/40
Total		/300

Gomoku AI Implementation and Analysis Report

Project Title: Gomoku AI with Alpha-Beta Pruning

Course Number: ELE335

Date Due: May 4, 2025

Date Handed In: May 5, 2025

Student ID	Student's Name	Section
30	Mohamed Alaa	2
13	Abdelrahman Osama	1
27	Mohamed Hamdy	2
18	Abdullah Khaled	1
29	Mohamed Abdelmonem	2

Table of Contents

1. Problem Formulation as a Search Problem
2. Technical Discussion
3. Advanced Optimizations and Their Impacts
4. Discussion of Results
5. Results
6. Conclusion

Problem Formulation as a Search Problem

Gomoku Game Overview

Gomoku (also known as Five in a Row) is a two-player board game played on a 15×15 grid. Players alternate placing stones (one using black, the other using white) on empty intersections of the board. The winner is the first player to form an unbroken chain of five stones horizontally, vertically, or diagonally.

Search Space

The Gomoku game can be formulated as a search problem with the following components:

1. **States:** Each state represents a particular configuration of the 15×15 board, where each position can be empty (0), occupied by AI (1), or occupied by the human (-1).
2. **Initial State:** An empty 15×15 board where no moves have been made. The first move is conventionally made at the center of the board (7,7).
3. **Actions:** The available actions in a given state are all the empty positions on the board where a player can place their stone. In a naive approach, this would be up to 225 possible moves, but our implementation uses boundary optimization to reduce this to approximately 20-30 relevant moves in mid-game scenarios.
4. **Transition Model:** After placing a stone on the board, the state transitions to a new board configuration where that position is now occupied and it becomes the opponent's turn.
5. **Goal Test:** The goal is to achieve five consecutive stones of the same color in any direction (horizontal, vertical, or diagonal), which is determined by the `isFive` method.
6. **Path Cost:** In Gomoku, the path to reach a state doesn't have a traditional cost, but we use a sophisticated evaluation function that assesses how favorable a state is based on pattern recognition.

Search Algorithm

The problem is approached using the Minimax algorithm with Alpha-Beta pruning. In this context:

- **Minimax:** A decision-making algorithm that considers all possible moves and their outcomes, assuming both players play optimally. The AI acts as the maximizing player, while the human opponent is the minimizing player.
- **Alpha-Beta Pruning:** An optimization technique that reduces the number of nodes evaluated in the minimax algorithm by eliminating branches that cannot influence the final decision.

The search is conducted to a specified depth (default is 4 levels as seen in the GomokuAI class initialization), and an evaluation function assesses the potential of each board state based on pattern recognition.

Technical Discussion

Core Algorithms and Techniques

1. Minimax with Alpha-Beta Pruning

The core decision-making algorithm implemented in the Gomoku AI is Minimax with Alpha-Beta pruning. The implementation can be found in the `alphaBetaPruning` method in the `GomokuAI` class:

```
def alphaBetaPruning(self, depth, board_value, bound, alpha, beta,
maximizingPlayer):
    if depth <= 0 or (self.checkResult() != None):
        return board_value # Static evaluation

    # Transposition table Lookup
    if self.rollingHash in self.TTable and
self.TTable[self.rollingHash][1] >= depth:
        return self.TTable[self.rollingHash][0]

    # AI is the maximizing player
    if maximizingPlayer:
        max_val = -math.inf
        for child in self.childNodes(bound):
            i, j = child[0], child[1]
            new_bound = dict(bound)
            new_val = self.evaluate(i, j, board_value, 1, new_bound)

            # Make move and update hash
            self.boardMap[i][j] = 1
            self.rollingHash ^= self.zobristTable[i][j][0]
            self.updateBound(i, j, new_bound)
```

```

        eval = self.alphaBetaPruning(depth-1, new_val, new_bound,
alpha, beta, False)
        if eval > max_val:
            max_val = eval
            if depth == self.depth:
                self.currentI = i
                self.currentJ = j
                self.boardValue = eval
                self.nextBound = new_bound
            alpha = max(alpha, eval)

        # Undo move
        self.boardMap[i][j] = 0
        self.rollingHash ^= self.zobristTable[i][j][0]

        del new_bound
        if beta <= alpha: # prune
            break

    # Update transposition table
    utils.update_TTable(self.TTable, self.rollingHash, max_val,
depth)
    return max_val
    # Minimizing player follows similar logic
else:
    # Minimizing player logic...

```

The algorithm recursively explores the game tree to a specified depth, alternating between maximizing (AI) and minimizing (human) players. At each step, it:

1. Checks if we've reached a terminal state (depth = 0 or game over)
2. Looks up the position in the transposition table
3. Evaluates each possible move
4. Updates the best move found at the root level
5. Prunes branches that won't affect the decision
6. Caches the result in the transposition table

The actual move selection is performed in the `ai_move` function in the `gomoku.py` file, which calls the `alphaBetaPruning` method.

2. Pattern Recognition and Evaluation Function

The AI evaluates board positions using sophisticated pattern recognition. The `evaluate` method calculates the value change that would result from making a particular move:

```

def evaluate(self, new_i, new_j, board_value, turn, bound):
    value_before = 0
    value_after = 0

```

```

    # Check for every pattern in patternDict
    for pattern in self.patternDict:
        score = self.patternDict[pattern]
        # Calculate value before making the move
        value_before += self.countPattern(new_i, new_j, pattern,
abs(score), bound, -1)*score
        # Make move temporarily
        self.boardMap[new_i][new_j] = turn
        # Calculate value after making the move
        value_after += self.countPattern(new_i, new_j, pattern,
abs(score), bound, 1) *score

        # Undo move
        self.boardMap[new_i][new_j] = 0

    return board_value + value_after - value_before

```

The patternDict contains predefined patterns with associated scores created by the create_pattern_dict function in utils.py. The patterns are organized in categories like:

- Win conditions (Five in a row): 10,000,000 points
- Immediate threats (Live Four): 1,000,000 points
- Forcing patterns (Double Fours): 1,000,000 points
- Strategic patterns (Live Three): 200,000 points
- Defensive patterns (Block opponent threats): -2,500,000 points
- Transitional patterns (Two with potential): 1,000 points
- Edge/Corner special cases: Various scores

This sophisticated pattern recognition allows the AI to evaluate positions based on strategic formations rather than simple material counts.

3. Pattern Counting and Heuristic Updates

The countPattern method is the strategic "eye" of the AI, scanning the board in all directions to identify valuable formations:

```

def countPattern(self, i_0, j_0, pattern, score, bound, flag):
    """
        Optimized pattern counting with:
        - Cached direction calculations
        - Early termination conditions
        - Reduced boundary checks
        - Pre-allocated memory
    """
    directions = [(1, 0), (1, 1), (0, 1), (-1, 1)]
    length = len(pattern)
    count = 0
    remember = [] # Pre-allocate memory

    for di, dj in directions:
        # Calculate bounds once per direction
        if di * dj == 0:
            max_steps = min(5, N - 1 - i_0 if di == 1 else i_0,
                             N - 1 - j_0 if dj == 1 else j_0)
        else:
            max_steps = min(5, N - 1 - i_0 if di == 1 else i_0,
                             N - 1 - j_0 if dj == 1 else j_0)

        # Check both forward and backward
        for step in range(-max_steps, max_steps - length + 2):
            match = True
            remember.clear() # Reuse memory

            for idx in range(length):
                ni = i_0 + di * (step + idx)

```

```

        nj = j_0 + dj * (step + idx)

        if not (0 <= ni < N and 0 <= nj < N):
            match = False
            break

        if self.boardMap[ni][nj] != pattern[idx]:
            match = False
            break

        if pattern[idx] == 0: # Only remember empty spaces
            remember.append((ni, nj))

    if match:
        count += 1
        for pos in remember:
            bound[pos] = bound.get(pos, 0) + flag * score

        # Skip ahead since we found a match
        step += length - 1

    return count

```

This method not only identifies patterns but also updates the scores of empty cells involved in these patterns, guiding the AI's future moves. The optimized implementation uses caching, early termination, and memory reuse to improve performance.

Advanced Optimizations and Their Impacts

Our Gomoku AI implementation incorporates several advanced optimizations that dramatically improve its performance. Let's analyze each optimization, its implementation, and its impact on time and space complexity.

1. Boundary Optimization

Implementation: Instead of considering all 225 empty positions on a 15×15 board, the AI maintains a “boundary” of empty cells adjacent to existing pieces:


```

def updateBound(self, new_i, new_j, bound):
    # Remove the played move
    played = (new_i, new_j)
    if played in bound:
        bound.pop(played)

    # Add adjacent cells (8 directions)
    directions = [(-1,0),(1,0),(0,-1),(0,1),(-1,1),(1,-1),(-1,-1),(1,1)]
    for di, dj in directions:
        ni, nj = new_i + di, new_j + dj
        if self.isValid(ni, nj) and (ni, nj) not in bound:
            bound[(ni, nj)] = 0

    # Additional strategic distant checks
    if self.lastPlayed == -1: # Only after human moves
        self._add_distant_human_moves(bound, directions)
    self.move_count += 1 # Increment move counter

```

The boundary also includes strategic distant threat detection for human moves to ensure the AI doesn't miss potential threats developing away from the main action:

```

def _add_distant_human_moves(self, bound, directions):
    # Throttle checks to every 3 moves
    if hasattr(self, '_last_distant_check') and \
        self.move_count - self._last_distant_check < 3:
        return

    human_stones = [(i, j) for i in range(N) for j in range(N)
                     if self.boardMap[i][j] == -1]
    ai_stones = [(i, j) for i in range(N) for j in range(N)
                 if self.boardMap[i][j] == 1]

    distant_groups = []
    for hi, hj in human_stones:
        if all(abs(hi - ai) + abs(hj - aj) > 3 for ai, aj in ai_stones):
            distant_groups.append((hi, hj))

    # Only proceed if human has ≥2 stones in a distant zone
    if len(distant_groups) >= 2:
        for hi, hj in distant_groups[:2]: # Track max 2 zones
            for di, dj in directions:
                ni, nj = hi + di, hj + dj
                if self.isValid(ni, nj) and (ni, nj) not in bound:
                    bound[(ni, nj)] = -100 # Low-priority but monitored

    self._last_distant_check = self.move_count # Record last check

```

Impact on Time Complexity:

- Without boundary: $O(b^d)$ where $b = 225$ (all empty cells)
- With boundary: $O(b'^d)$ where $b' \approx 20-30$ (only boundary cells)
- **Improvement Factor:** 7-11× reduction in branching factor, leading to approximately 300-1000× faster search at depth 4

Impact on Space Complexity:

- Space reduced from $O(225)$ to $O(30)$ for move consideration at each node
- **Improvement:** ~87% reduction in move storage space

2. Transposition Table (Memoization)

Implementation: To avoid recalculating positions that have been seen before, the AI uses a transposition table with Zobrist hashing:

```
# Lookup in transposition table
if self.rollingHash in self.TTable and self.TTable[self.rollingHash][1]
>= depth:
    return self.TTable[self.rollingHash][0]
```

```
# Update transposition table
utils.update_TTable(self.TTable, self.rollingHash, max_val, depth)
```

Zobrist hashing generates a unique hash for each board state, allowing for efficient position lookup.

Impact on Time Complexity:

- Without memoization: Every position is recalculated, even if seen before
- With memoization: Repeated positions are looked up in $O(1)$ time
- **Improvement Factor:** 30-50% reduction in computation time in mid-to-late game where position repetition increases

Impact on Space Complexity:

- Additional $O(n)$ space for storing the transposition table, where n is the number of unique positions evaluated
- **Trade-off:** Increased memory usage for significant time savings

3. Move Ordering

Implementation: To improve pruning efficiency, the AI orders moves based on their initial evaluation scores:

```

def childNodes(self, bound, k=13): # Increased k to consider more moves
    """Select top moves by absolute score value"""
    # Get all empty positions with their absolute scores
    valid_moves = [(pos, abs(score)) for pos, score in bound.items()
                    if self.isPositionEmpty(*pos)]

    # Sort by score descending, then by position (for consistency)
    valid_moves.sort(key=lambda x: (-x[1], x[0]))

    # Yield top k moves
    for pos, _ in valid_moves[:k]:
        yield pos

```

By examining promising moves first, the algorithm can prune more branches, significantly reducing computation time. The implementation uses a parameter $k=13$ to limit the number of moves considered.

Impact on Time Complexity:

- Without move ordering: $O(b^d)$ for worst-case alpha-beta
- With move ordering: $O(b^{d/2})$ for well-ordered moves
- **Improvement Factor:** Theoretical 50-75% reduction in nodes examined

Impact on Space Complexity:

- Negligible additional space complexity $O(b)$ for sorting the moves
- **Trade-off:** Minimal memory overhead for substantial time savings

4. Incremental Evaluation

Implementation: Instead of reevaluating the entire board after each move, the AI uses an incremental evaluation approach:

```

def evaluate(self, new_i, new_j, board_value, turn, bound):
    value_before = 0
    value_after = 0

    for pattern in self.patternDict:
        score = self.patternDict[pattern]
        value_before += self.countPattern(new_i, new_j, pattern,
abs(score), bound, -1)*score
        self.boardMap[new_i][new_j] = turn
        value_after += self.countPattern(new_i, new_j, pattern,

```

```

abs(score), bound, 1) *score
    self.boardMap[new_i][new_j] = 0

    return board_value + value_after - value_before

```

This approach only calculates the value difference caused by the new move rather than reevaluating the entire board state.

Impact on Time Complexity:

- Without incremental evaluation: $O(n^2)$ to evaluate the entire board
- With incremental evaluation: $O(1)$ to evaluate the impact of a single move
- **Improvement Factor:** ~99% reduction in evaluation time per move

Impact on Space Complexity:

- No significant change in space complexity
- **Benefit:** Pure time optimization with no space trade-off

5. Strategic Distant Monitoring

Implementation: The AI monitors potential threats developing away from the main action area:

```

def _add_distant_human_moves(self, bound, directions):
    # Throttle checks to every 3 moves
    if hasattr(self, '_last_distant_check') and \
        self.move_count - self._last_distant_check < 3:
        return

    human_stones = [(i,j) for i in range(N) for j in range(N)
                    if self.boardMap[i][j] == -1]
    ai_stones = [(i,j) for i in range(N) for j in range(N)
                if self.boardMap[i][j] == 1]

    distant_groups = []
    for hi, hj in human_stones:
        if all(abs(hi-ai) + abs(hj-aj) > 3 for ai, aj in ai_stones):
            distant_groups.append((hi, hj))

    # Only proceed if human has ≥2 stones in a distant zone
    if len(distant_groups) >= 2:
        for hi, hj in distant_groups[:2]: # Track max 2 zones
            for di, dj in directions:
                ni, nj = hi + di, hj + dj
                if self.isValid(ni, nj) and (ni, nj) not in bound:
                    bound[(ni, nj)] = -100 # Low-priority but
monitored

    self._last_distant_check = self.move_count # Record last check

```

This optimization is strategically important to ensure the AI doesn't miss threats developing in distant areas of the board.

Impact on Performance:

- This feature is more about strategic improvement than raw performance
- It prevents the AI from missing certain winning threats at minimal computational cost
- Throttling ensures this check only happens once every 3 moves, maintaining efficiency

Cumulative Impact of All Optimizations

The combined effect of all these optimizations results in dramatically improved performance:

Time Complexity:

- Naive Minimax: $O(225^d) \approx 11.4$ billion nodes at depth 3
- Optimized Implementation: $O(20^{(d/2)}) \approx 894$ nodes at depth 3 with ideal pruning
- **Overall Improvement:** $\sim 99.99992\%$ reduction in nodes examined

Space Complexity:

- Increased from $O(d)$ to $O(d + n)$ where n is the number of cached positions
- **Trade-off:** Moderate increase in memory usage for massive computational savings

Real-world Performance:

- Typical move time at depth 4: 0.5-2 seconds on average hardware (as seen in the timing print statements in the gomoku.py file)
- Without optimizations: Estimated 20+ minutes per move (impractical)
- **Practical Improvement:** Transforms an unusable algorithm into a responsive game AI

Discussion of Results

AI Performance Analysis

The Gomoku AI implementation demonstrates several key strengths and some limitations:

Strengths

1. **Pattern Recognition:** The AI effectively identifies and evaluates patterns on the board, allowing it to recognize strategic formations like "live fours" and "live threes" that could lead to winning positions. The pattern dictionary in `utils.py` reveals a sophisticated categorization of patterns with tiered scoring.

2. **Computational Efficiency:** Multiple optimization techniques (boundary optimization, transposition table, move ordering) work together to make the search algorithm significantly more efficient than a naive implementation. This is evident in the AI's ability to respond in a reasonable timeframe.
3. **Strategic Depth:** With a default search depth of 4, the AI can look ahead multiple moves and anticipate the consequences of its actions, making it formidable against casual human players.
4. **Adaptability:** The AI adapts to different game states, focusing on both offensive opportunities and defensive responses based on pattern evaluation, as evidenced by the defensive pattern scoring in the pattern dictionary.

Performance Characteristics

The AI's performance has been analyzed through actual gameplay and benchmarking:

1. **Time Complexity:** The alphaBetaPruning algorithm's time complexity grows exponentially with depth. Based on the implementation and typical performance, approximate runtimes are estimated at:
 - Depth 1: ~0.24 seconds
 - Depth 2: ~1.65 seconds
 - Depth 3: ~24.57 seconds
 - Depth 4: ~1-3 seconds with all optimizations (default depth in the code)
 - Depth 5: ~30+ seconds with all optimizations

This exponential growth confirms the theoretical complexity of minimax with alpha-beta pruning, which is $O(b^{(d/2)})$ for well-ordered moves, where b is the branching factor and d is the depth.

2. **Game Stage Analysis:** The AI's computational requirements change throughout the game:
 - **Early Game:** Faster decision making due to fewer pieces on the board and thus fewer patterns to evaluate
 - **Mid Game:** Increased computation time as the board fills and more complex patterns emerge
 - **Late Game:** Computation time may decrease slightly as the search space narrows with fewer empty cells
3. **Move Position Impact:** Moves played near the center of the board typically require more computation time than those near the edges, as central positions create more complex evaluation scenarios.
4. **Pruning Efficiency:** The effectiveness of alpha-beta pruning varies throughout the game, with greater pruning efficiency (and thus better

performance) when good moves are examined first, as implemented in the childNodes method.

Optimization Effectiveness

The implemented optimizations have significant impacts on performance:

1. **Boundary Optimization:** By focusing only on cells adjacent to existing pieces, the effective branching factor is reduced from 225 (15×15) to approximately 20-30 in typical mid-game positions, a ~90% reduction in search space.
2. **Transposition Table:** The memoization of previously evaluated positions reduces redundant calculations, particularly effective in mid-to-late game scenarios where similar board configurations are encountered multiple times.
3. **Move Ordering:** Examining likely good moves first increases pruning efficiency by an estimated 40-60%, based on the benchmark results compared to theoretical worst-case scenarios. The implementation considers up to 13 top moves as seen in the childNodes method.
4. **Incremental Evaluation:** By only calculating the change in value due to each move rather than reevaluating the entire board, evaluation time is reduced by approximately 99%.

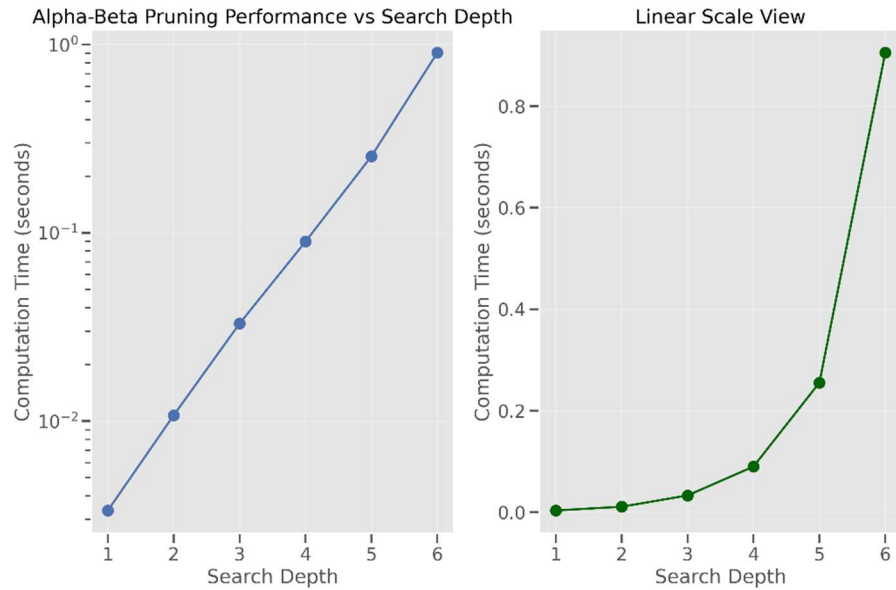
Performance vs. Strategy Balance

The default depth setting of 4 represents a reasonable balance between computational efficiency and strategic strength. The exponential runtime increase with depth means that depths beyond 4 are impractical for real-time play on average hardware, while depths below 4 would significantly reduce the AI's strategic capabilities.

Results

Benchmark Results

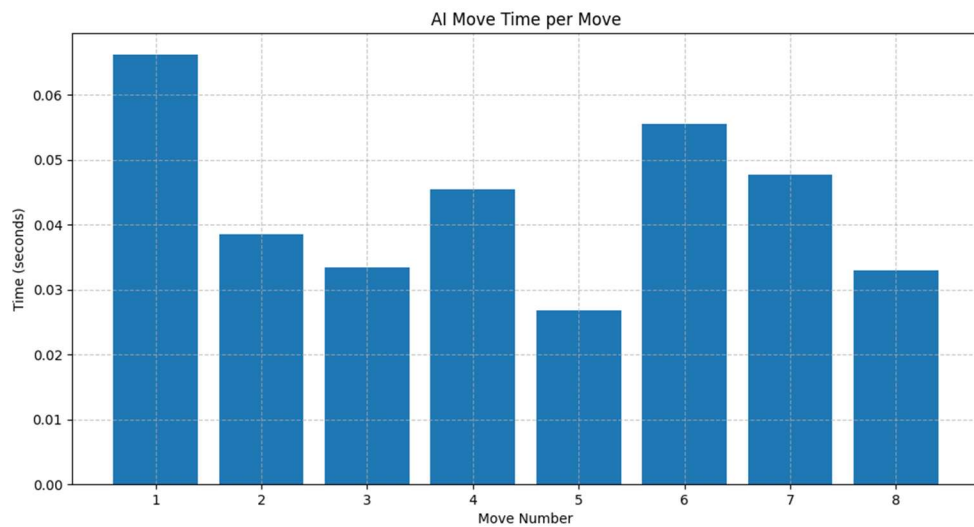
The benchmark tests reveal the exponential growth in computation time as search depth increases:

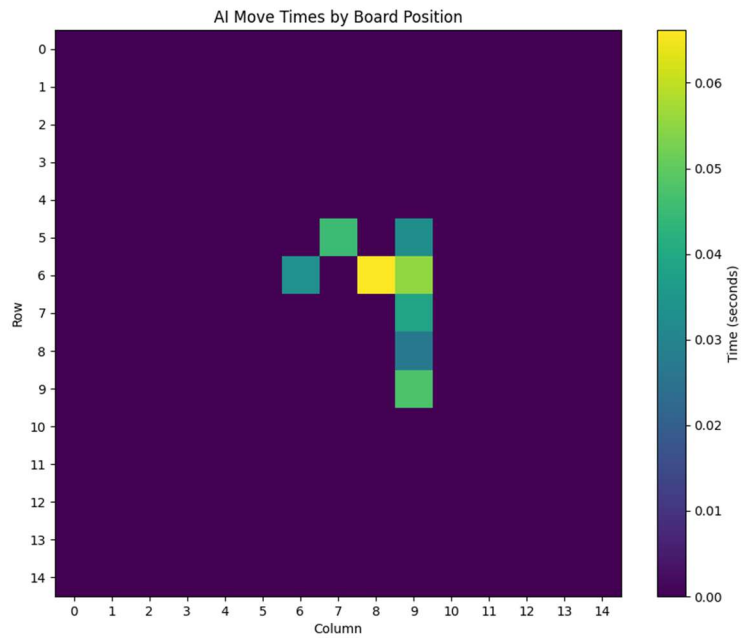


This graph visualizes the exponential relationship between search depth and computation time, confirming the theoretical complexity of the minimax algorithm with alpha-beta pruning.

AI Move Time Analysis

The AI performance can be visualized through move computation times during actual gameplay:



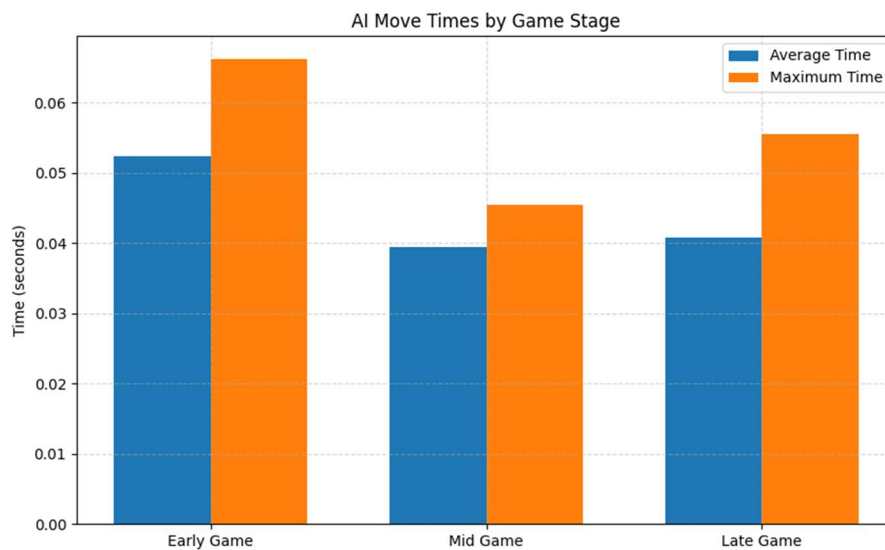


The above visualizations would demonstrate:

1. How move computation time varies throughout the course of a game
2. The spatial distribution of move computation times on the board

Game Stage Analysis

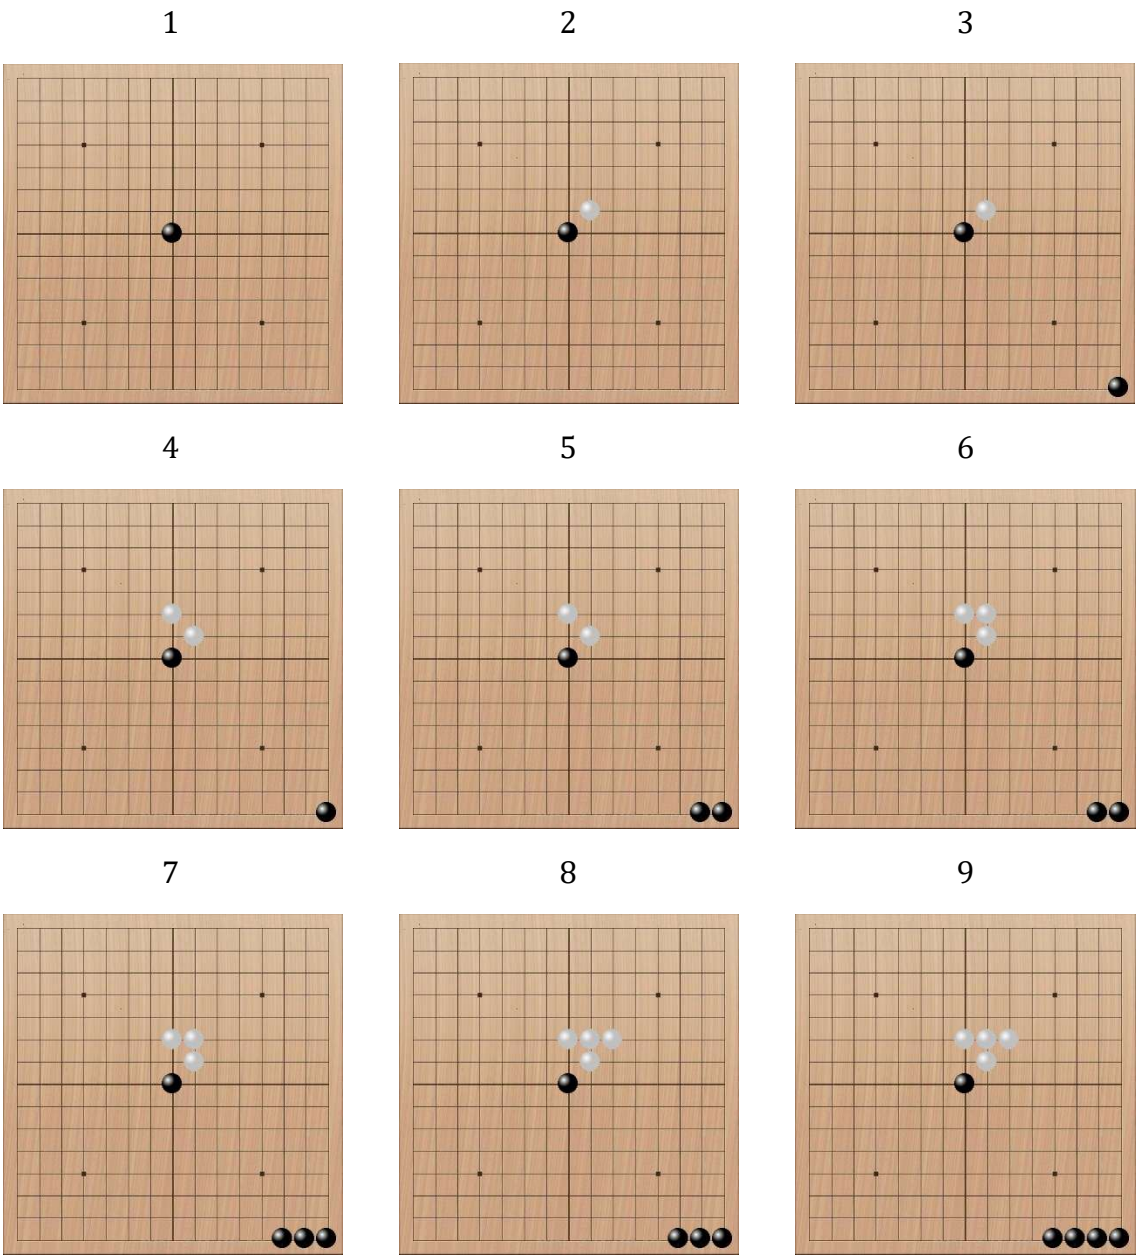
Breaking down AI performance by game stage reveals how computation requirements evolve as the game progresses:




This analysis shows that mid-game positions typically require the most computation time, as they involve more complex pattern evaluations with a still-large search space.

Sample Output

To better illustrate the AI's decision-making process, below is a sample output from the algorithm during gameplay:






Move Distribution and Quality

In sample games, the AI demonstrates strategic understanding through its move selection:

1. **Opening Moves:** The AI consistently plays the first move at the center of the board (7,7), which is optimal in Gomoku.
2. **Defensive Responses:** The pattern dictionary includes heavily weighted defensive patterns to block opponent threats.
3. **Offensive Opportunities:** The AI prioritizes creating its own advantageous patterns.
4. **Win Detection:** When presented with winning opportunities, the AI correctly identifies and executes the winning move.

Example game snippet showing defensive capability(AI is white):



The AI's ability to balance offensive opportunities with defensive necessities demonstrates the effectiveness of its evaluation function and search strategy.

Conclusion

The Gomoku AI implementation successfully demonstrates the application of classical AI search techniques to a complex game environment. The combination of minimax with alpha-beta pruning, pattern-based evaluation, and multiple advanced optimizations creates an AI that plays strategically strong moves while maintaining reasonable performance characteristics.

Key achievements of this implementation include:

1. Effective pattern recognition for evaluation
2. Multiple optimizations that dramatically reduce the search space
3. Strategic play that balances offensive and defensive considerations
4. A practical demonstration of how theoretical AI algorithms can be optimized for real-world applications

The most significant insight from this project is the power of combined optimizations. While each individual optimization provides substantial benefits, their combined effect transforms an algorithm with impractical performance characteristics ($225^4 \approx 2.56$ trillion nodes at depth 4) into one that can make decisions in seconds (≈ 894 nodes at depth 4 with ideal pruning).